

Berkeley Packet Filter

Linux kernel TLV meetup, 28.02.2016

kerneltlv.com

Kfir Gollan

Agenda

- What is the Berkeley Packet Filter?
- Writing BPF filters
- Debugging BPF
- Using BPF in user-space applications
- Advanced features of BPF

What is the Berkeley Packet Filter?

- BPF at its base is a way to perform fast packet filtering at the kernel level.
 - Filters are defined by user space
 - Filters are executed in the kernel
- Invented by Steven McCanne and Van Jacobson in 1990. First publication at December 1992.
- Support for BPF in Linux was added by Jay Schulist for the 2.5 development kernel.
 - Many features were added later on.
For example JIT for BPF was added for the 3.0 kernel (2011)

Why do we need BPF?

- We want to be able to filter packets at the kernel level
 - Discard irrelevant packets at the kernel without copying them to user space.
 - The performance gains when using promiscuous mode are substantial.
- We want to change filters dynamically without recompiling the kernel or using a custom kernel module.
- We want the filters to be architecture independent.

BPF in wireshark

The Wireshark Network Analyzer [Wireshark 1.8.2]

File Edit View Go Capture Analyze Statistics Telephony Tools Internals Help

Edit Interface Settings

Capture

Interface: eth0

IP address: 192.168.10.128
fe80::a00:27ff:fe1e:ddbd

Link-layer header type: Ethernet Buffer size: 1 megabyte(s)

☒ Capture packets in promiscuous mode ☐ Capture packets in monitor mode

☐ Limit each packet to 65535 bytes

Capture Filter: src port 80 Compile BPF

Help Cancel OK

Linux netfilter log (NFLOG) interface: nflog

Pseudo-device that captures on all interfaces: any

Capture Options

Start a capture with detailed options

BPF Filter

(000) ldh [12]
(001) jeq #0x86dd jt 2 jf 8
(002) ldb [20]
(003) jeq #0x84 jt 6 jf 4
(004) jeq #0x6 jt 6 jf 5
(005) jeq #0x11 jt 6 jf 19
(006) ldh [54]
(007) jeq #0x50 jt 18 jf 19
(008) jeq #0x800 jt 9 jf 19
(009) ldb [23]
(010) jeq #0x84 jt 13 jf 11
(011) jeq #0x6 jt 13 jf 12
(012) jeq #0x11 jt 13 jf 19
(013) ldh [20]
(014) jset #0x1fff jt 19 jf 15
(015) ldx 4*([14]&0xf)
(016) ldh [x + 14]
(017) jeq #0x50 jt 18 jf 19
(018) ret #65535
(019) ret #0

OK

Writing BPF

BPF basics

- BPF defines a set of operations that can be performed on the filtered packet. Each operation gets its own opcode.
- BPF was designed to be protocol independent, as such it treats the packets as raw buffers. It is up to the filter writer to parse the needed packet headers.
- BPF is based on three building blocks:
 - A: A 32 bit wide accumulator
 - X: A 32 bit wide index register
 - M[]: 16 x 32 bit wide misc registers aka “scratch memory”

BPF instruction types

- **LOAD:** copy a value into the accumulator or index register.
- **STORE:** copy either the accumulator or index register to the scratch memory.
- **ALU:** perform arithmetic or logic operation on the accumulator register.
- **BRANCH:** alter the flow of control.
- **RETURN:** terminate the filter and indicate what portion of the packet to save.
- **MISC:** various operations that doesn't match the other types.

BPF instruction format

- Each instruction is represented by 64 bits.
 - opcode 16 bits, indicates the instruction type.
 - jt 8 bits, offset of the next instruction for true case (“if” block) – Jump True.
 - jf 8 bits, offset of the next instruction for false case (“else” block) – Jump False.
 - k 32 bits, generic data, used for various purposes.

opcode: 16	jt: 8	jf: 8
k: 32		

BPF instruction set

Instruction	Description	Instruction	Description
ld	Load word into A	add	$A + \langle x \rangle$
ldh	Load half-word into A	sub	$A - \langle x \rangle$
ldb	Load byte into A	mul	$A * \langle x \rangle$
ldx	Load word into X	div	$A / \langle x \rangle$
ldxb	Load byte into X	and	$A \& \langle x \rangle$
st	Store A into M[]	or	$A \langle x \rangle$
stx	Store X into M[]	xor	$A \wedge \langle x \rangle$
jmp	Jump to label	lsh	$A \ll \langle x \rangle$
jeq	Jump on $k == A$	rsh	$A \gg \langle x \rangle$
jgt	Jump on $k > A$	ret	Return
jge	Jump on $k \geq A$	tax	Copy A into X
jset	Jump on $k \& A$	txa	Copy X into A

BPF addressing modes

Mode	Description
#k	The literal value stored in k.
#len	The length of the packet.
M [k]	The word at offset k in the scratch memory store.
[k]	The byte, half-word or word at byte offset k in the packet.
[x + k]	The byte, half-word or word at byte offset x+k in the packet.
L	Jump to label L
#k, lt. lf	Jump to lt if true, otherwise jump to l f
x	The index register
4 * ([k] & 0xf)	Four times the value of the low four bits of the byte at the offset k in the packet

Example headers

- MAC header

Destination MAC	Source MAC	Type
6 bytes	6 bytes	2 bytes

- IP header

4-bit	8-bit	16-bit	32-bit	
Ver.	Header Length	Type of Service	Total Length	
Identification			Flags	Offset
Time To Live	Protocol		Checksum	
Source Address				
Destination Address				
Options and Padding				

Example 1: All IP packets

- To catch all the IP packets over MAC we need to check the type field in the MAC header.

	ldh	[12]
	jeq	#ETHERTYPE_IP, L1, L2
L1:	ret	#-1
L2:	ret	#0

- Load half-word (2 bytes) from offset 12 of the packet (the type field) into A register.
- Check if A register is equal to #ETHERTYPE_IP
 - If true -> return #-1
 - If false -> return 0

Example 2: IP not from 128.3.112.X

	ldh	[12]	; A <= ether.type
	jeq	#ETHERTYPE_IP, L1, L3	; A == #ETHERTYPE_IP ?
L1:	ld	[26]	; A = ip.src
	and	#0xffffffff	; A = A & 0xffffffff
	jeq	#0x80037000, L3, L2	; A == 128.3.112.0
L2:	ret	#-1	
L3:	ret	#0	

- Check if the packet type is IP
- “Remove” the lower byte of the src IP
- Check if the src IP matches 128.3.112.X

Debugging BPF

bpf-asm

- A utility used to create bpf binary code (bpf “assembly” compiler”).
 - Part of the mainline kernel. `tools/net/bpf_asm.c`
- Supports two output formats
 - c style output

```
{ 0x28, 0, 0,    0x00000000c  },  
{ 0x15, 0, 1,    0x000000800  },  
{ 0x06, 0, 0,    0xffffffff  },  
{ 0x06, 0, 0,    00000000000  },
```
 - raw output

```
4,40 0 0 12,21 0 1 2048,6 0 0 4294967295,6 0 0 0,
```


bpf-dbg

- A utility used to debug bpf filters
 - Part of the mainline kernel. `tools/net/bpf_dbg.c`
- Main features
 - pcap files as input for filters.
 - bpf-asm raw output format for bpf filter definition.
 - single-stepping through filters
 - breakpoints
 - internal status (A,X,M, PC)
 - disassemble raw bpf to bpf-asm

Debugging demo

Using BPF in user-space

linux/filter.h – sock_filter

```
struct sock_filter {                                /* Filter block */
    __u16    code;                                  /* Actual filter code */
    __u8     jt;                                    /* Jump true */
    __u8     jf;                                    /* Jump false */
    __u32    k;                                     /* Generic multiuse field */
};
```

- Filter block is in fact a single BPF instruction.
- Used to pass a filter specifications to the kernel.

```
struct sock_filter code[] = {
    { 0x28, 0, 0, 0x00000000c },
    { 0x15, 0, 8, 0x0000086dd },
    { 0x30, 0, 0, 0x000000014 },
    ...
};
```


linux/filter.h – sock_fprog

```
struct sock_fprog {                                /* Required for SO_ATTACH_FILTER. */  
    unsigned short len;                            /* Number of filter blocks */  
    struct sock_filter *filter;                    /* Filter blocks list */  
};
```

- A parameter for setsockopt that allows to attach a filter to a socket.

```
struct sock_fprog bpf = {...};  
setsockopt(int fd,  
           SOL_SOCKET,  
           SO_ATTACH_FILTER,  
           &bpf,  
           sizeof(bpf));
```

setsockopt flags

- `SO_ATTACH_FILTER`
Attach a BPF filter to a socket.
Note: only a single filter can be attached at a given time.
- `SO_DETACH_FILTER`
Remove the currently attached filter from the socket.
- `SO_LOCK_FILTER`
Lock a filter on a socket. This is useful for setting a filter and then dropping privileges.
 - Example: create a raw socket, apply a filter, lock it, drop `CAP_NET_RAW`

Choosing socket() flags correctly

- Choosing the correct flags to the socket is critical for making the filter work properly.
 - The filtered buffer will start in the wanted location in the net stack.
- Selecting the socket domain
 - AF_PACKET – filtering at L2 (e.g ethernet)
 - AF_INET – IPv4 filtering
 - AF_INET6 – IPv6 filtering
- Selecting the socket type
 - SOCK_RAW – raw filtering, no headers are handled by the kernel
 - SOCK_STREAM/SOCK_DGRAM etc – headers are handled by the kernel

PCAP

- libpcap – Packet CAPture library
- Provides an easy to use api for packet filtering.
- Supports a high level filtering format which is converted to BPF.
 - pcap_compile – create a bpf filter
 - pcap_setfilter – attach a filter

```
ether dst [mac address]  
dst net [ip address]  
dst portrange [port1]-[port2]
```

- Look at man 7 pcap-filter for a detailed description.

tcpdump

- user-space packet sniffing program.
- Uses bpf for kernel level filtering (based on pcap)
 - **Dump the generated bpf filter**
 - -d bpf asm format
 - -dd c format
 - -ddd bpf raw format

```
$ sudo tcpdump -d "ip and udp"
(000) ldh     [12]
(001) jeq     #0x800            jt 2   jf 5
(002) ldb     [23]
(003) jeq     #0x11            jt 4   jf 5
(004) ret     #65535
(005) ret     #0
```

BPF - Advanced features

JIT

- A just-in-time BPF instruction translation.
 - Note that the translation is performed when attaching the filter via `bpf_jit_compile(..)`.
- BPF instructions are mapped directly to architecture depended instructions.
 - BPF registers are mapped to machine physical registers
- Provides a performance gain
 - About 50ns per packet for simple filters (E5540 @ 2.53GHz).
The more complex the filter the better performance gains from JIT.
- Supported on x86, x86-64, powerpc. arm and more.

BPF verifier

- Each BPF filter is verified before attaching it to a socket.
 - This is critical, the filters come from userspace!
- The following rules are enforced
 - The filter must not contain references or jumps that are out of range.
 - The filter must contain only valid BPF opcodes.
 - The filter must end with RET opcode.
 - All jumps are forward – loops are not allowed!
- The verification is implemented at `sk_chk_filter` function in `net/core/filter.c` (until kernel 3.16), modified after adding `seccomp`.

BPF extensions

- The Linux kernel also has a couple of BPF extensions that are used along with the class of load instructions.
- The extensions are "overloading" the k argument with a negative offset + a particular extension offset.
- The result of such BPF extensions are loaded into A.

Instruction	Description	Instruction	Description
len	skb->len	queue	skb->queue_mapping
proto	skb->protocol	rxhash	skb->hash
type	skb->pkt_type	rand	Prandom_u32()
poff	Payload start offset	cpu	Executing cpu id
ifidx	skb->dev->ifindex	nla	Netlink attributes
mark	skb->mark	hatype	skb->dev->type

eBPF

- extended BPF is an internal mechanism that can be used only in kernel context (not from userspace!)
- eBPF adds a set of new features:
 - Increased number of registers 10 instead of 2.
 - Register width increased to 64 bit
 - Conditional jf/jt targets replaced with jt/fall-through
 - bpf_call instruction and register passing convention for zero overhead calls from/to other kernel functions.
- Originally designed to be a “restricted C” language that will be architecture independent and JITed in kernel context.
- Eventually used mostly for kernel tracing.

BPF – not only for networking

SecComp

- SECure COMPuting, or seccomp, is a security mechanism available in the linux kernel.
- It applies BPF filtering to syscalls
 - filter the syscall number & its parameters
- It can be used to limit the available syscalls
 - For example strict mode allows only read, write, _exit and sigreturn
- Uses the BPF filters, the filtered buffers are different.
- Look at man 2 seccomp for more details.

Questions?
